

# RAG Project Planning

## Level 1: Mini RAG Project (1–2 days)

**Project Idea:** *"Ask Questions About My Notes/PDF"*

- Take a small PDF (resume, research paper, course notes).
- Implement RAG so you can ask questions like:
  - "What's the summary of this PDF?"
  - "What is mentioned about project X in the resume?"
- Stack: LangChain + FAISS + OpenAI API (or free Llama 3 on HF).

 LinkedIn Post Idea:

*"Built my first RAG chatbot that answers questions from my own PDF using LangChain + FAISS. No more scrolling through 50 pages 🚀"*

---

## Level 2: Medium Project (1–2 weeks)

**Project Idea:** *"Company Policy Chatbot"*

- Collect public HR/leave/IT policy PDFs (or dummy docs).
- Build a chatbot where employees can ask:
  - "How many casual leaves do I get?"
  - "What is the WFH policy?"
- Add **UI with Streamlit/Gradio**.
- Deploy free on Hugging Face Spaces.

 LinkedIn Post Idea:

*"Tired of searching in PDFs? I built a chatbot that answers HR policy questions instantly. Uses Retrieval-Augmented Generation (RAG) with LangChain + FAISS + Streamlit."*

---

## 🟡 Level 3: Advanced Project (Iteration)

Take your HR bot (or PDF bot) and **make it better** step by step:

1. Add **hybrid search (BM25 + embeddings)** for accuracy.
2. Add **source citations** ("Answer based on Page 12 of Policy.pdf").
3. Add **chat history memory** (multi-turn conversation).
4. Add **evaluation** (use LlamaIndex evals or RAGAS to test hallucination rate).
5. Deploy on **Docker + AWS/GCP/Render** for production feel.

✅ LinkedIn Post Idea (iteration):

*"Upgraded my RAG chatbot! Now it not only answers HR policy questions but also cites exact sources and handles follow-up queries. Deployed on the cloud 🚀"*

---

## 🔴 Level 4: Portfolio-worthy Project (1–2 months)

**Project Idea:** "RAG-Powered Knowledge Hub"

- Upload multiple document types (PDFs, CSVs, websites, YouTube transcripts).
- User selects a data source → chatbot answers queries.
- Add **admin panel** to upload new documents (so knowledge updates dynamically).
- Use **Weaviate/Pinecone** for scalable retrieval.
- Experiment with **RAG 2 (per-token retrieval)** or **hybrid models**.

✅ LinkedIn Post Idea:

*"From a small PDF Q&A bot to a full-fledged RAG-powered Knowledge Hub — now supporting multiple file types, cloud deployment, and scalable vector DB. Learned tons about embeddings, retrieval, and hallucination handling along the way!"*

---

## 🔑 Strategy for LinkedIn

- Post each milestone (don't wait for the "perfect" project).
  - Show **screenshots/demos**.
  - Write a short "what I learned" section each time.
  - Tag relevant hashtags (#AI, #LangChain, #RAG, #GenerativeAI).
- 

👉 So the flow for you could be:

1. **Start tiny** → PDF Q&A bot.
2. **Iterate** → HR/Policy chatbot with UI.
3. **Enhance** → Add citations, hybrid search, memory.
4. **Scale** → Knowledge Hub with multi-source + deployment.